

mnist_mlp

December 3, 2021

```
[1]: '''Trains a simple deep NN on the MNIST dataset.  
  
Gets to 98.40% test accuracy after 20 epochs  
(there is *a lot* of margin for parameter tuning).  
2 seconds per epoch on a K520 GPU.  
'''  
  
from tensorflow import keras  
from tensorflow.keras.datasets import mnist  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense, Dropout  
from tensorflow.keras.optimizers import RMSprop  
  
batch_size = 128  
num_classes = 10  
epochs = 20  
  
# the data, split between train and test sets  
(x_train, y_train), (x_test, y_test) = mnist.load_data()  
  
x_train = x_train.reshape(60000, 784)  
x_test = x_test.reshape(10000, 784)  
x_train = x_train.astype('float32')  
x_test = x_test.astype('float32')  
x_train /= 255  
x_test /= 255  
print(x_train.shape[0], 'train samples')  
print(x_test.shape[0], 'test samples')  
  
# convert class vectors to binary class matrices  
y_train = keras.utils.to_categorical(y_train, num_classes)  
y_test = keras.utils.to_categorical(y_test, num_classes)  
  
model = Sequential()  
model.add(Dense(512, activation='relu', input_shape=(784,)))  
model.add(Dropout(0.2))  
model.add(Dense(512, activation='relu'))
```

```

model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer=RMSprop(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                   batch_size=batch_size,
                   epochs=epochs,
                   verbose=1,
                   validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

60000 train samples
10000 test samples
Model: "sequential"

| Layer (type) | Output Shape | Param # |
|---------------------|--------------|---------|
| dense (Dense) | (None, 512) | 401920 |
| dropout (Dropout) | (None, 512) | 0 |
| dense_1 (Dense) | (None, 512) | 262656 |
| dropout_1 (Dropout) | (None, 512) | 0 |
| dense_2 (Dense) | (None, 10) | 5130 |

Total params: 669,706
Trainable params: 669,706
Non-trainable params: 0

```

Epoch 1/20
469/469 [=====] - 5s 10ms/step - loss: 0.2450 -
accuracy: 0.9241 - val_loss: 0.1008 - val_accuracy: 0.9681
Epoch 2/20
469/469 [=====] - 4s 10ms/step - loss: 0.1014 -
accuracy: 0.9693 - val_loss: 0.0847 - val_accuracy: 0.9731
Epoch 3/20
469/469 [=====] - 5s 10ms/step - loss: 0.0761 -
accuracy: 0.9765 - val_loss: 0.0848 - val_accuracy: 0.9761

```

Epoch 4/20
469/469 [=====] - 5s 10ms/step - loss: 0.0597 - accuracy: 0.9823 - val_loss: 0.0681 - val_accuracy: 0.9807

Epoch 5/20
469/469 [=====] - 5s 10ms/step - loss: 0.0511 - accuracy: 0.9843 - val_loss: 0.0712 - val_accuracy: 0.9815

Epoch 6/20
469/469 [=====] - 5s 10ms/step - loss: 0.0436 - accuracy: 0.9869 - val_loss: 0.0756 - val_accuracy: 0.9819

Epoch 7/20
469/469 [=====] - 5s 10ms/step - loss: 0.0386 - accuracy: 0.9887 - val_loss: 0.0789 - val_accuracy: 0.9833

Epoch 8/20
469/469 [=====] - 4s 10ms/step - loss: 0.0335 - accuracy: 0.9897 - val_loss: 0.0871 - val_accuracy: 0.9821

Epoch 9/20
469/469 [=====] - 4s 9ms/step - loss: 0.0304 - accuracy: 0.9906 - val_loss: 0.0885 - val_accuracy: 0.9834

Epoch 10/20
469/469 [=====] - 4s 9ms/step - loss: 0.0329 - accuracy: 0.9912 - val_loss: 0.0797 - val_accuracy: 0.9836

Epoch 11/20
469/469 [=====] - 4s 10ms/step - loss: 0.0257 - accuracy: 0.9922 - val_loss: 0.0898 - val_accuracy: 0.9830

Epoch 12/20
469/469 [=====] - 5s 10ms/step - loss: 0.0259 - accuracy: 0.9925 - val_loss: 0.1038 - val_accuracy: 0.9825

Epoch 13/20
469/469 [=====] - 4s 9ms/step - loss: 0.0232 - accuracy: 0.9931 - val_loss: 0.0994 - val_accuracy: 0.9825

Epoch 14/20
469/469 [=====] - 4s 9ms/step - loss: 0.0222 - accuracy: 0.9938 - val_loss: 0.0995 - val_accuracy: 0.9849

Epoch 15/20
469/469 [=====] - 4s 9ms/step - loss: 0.0226 - accuracy: 0.9936 - val_loss: 0.1026 - val_accuracy: 0.9821

Epoch 16/20
469/469 [=====] - 4s 9ms/step - loss: 0.0207 - accuracy: 0.9941 - val_loss: 0.0991 - val_accuracy: 0.9850

Epoch 17/20
469/469 [=====] - 4s 9ms/step - loss: 0.0195 - accuracy: 0.9950 - val_loss: 0.1154 - val_accuracy: 0.9834

Epoch 18/20
469/469 [=====] - 4s 9ms/step - loss: 0.0183 - accuracy: 0.9951 - val_loss: 0.1172 - val_accuracy: 0.9845

Epoch 19/20
469/469 [=====] - 4s 9ms/step - loss: 0.0164 - accuracy: 0.9954 - val_loss: 0.1244 - val_accuracy: 0.9838

```
Epoch 20/20
469/469 [=====] - 4s 9ms/step - loss: 0.0169 -
accuracy: 0.9951 - val_loss: 0.1170 - val_accuracy: 0.9856
Test loss: 0.11704535782337189
Test accuracy: 0.9855999946594238
```

[]: